

# TLA+ and a Concurrent Hashmap Specification

Åsmund Aqissiaq Arild Kløvstad

Universitetet i Bergen

March 17, 2022

# Table of Contents

Introduction

The Hashmap

TLA+

The Formalization

Results

Lessons Learned and Conclusion

# Outline

1. What was the project?
2. The data structure
3. TLA+
4. The specification
5. Results and lessons learned

# The Project

- ▶ 10 pt optional Bachelor Thesis at UiT
- ▶ Spring semester 2020
- ▶ Supervised by Håvard D. Johansen
- ▶ Specifying and modelchecking a hash table in TLA+

# Problem Statement

- ▶ Formal methods
- ▶ Concurrent and distributed systems
- ▶ A novel hashmap [1]

*Shalev et. al.'s Hash Table can be formally verified for concurrent settings*

# Table of Contents

Introduction

The Hashmap

TLA+


The Formalization

Results


Lessons Learned and Conclusion

## The Idea


- ▶ Resizing is hard in concurrent settings
- ▶ Solution: move the buckets instead
- ▶ Split-ordering (sort by reverse binary representation)
- ▶ Dummy nodes




## Insertion with Splitting




(a) Initial state



(b) Table is expanded



(c) Dummy node for new bucket inserted



(d) Table entry points to bucket node



# Table of Contents

Introduction

The Hashmap

TLA+

The Formalization

Results

Lessons Learned and Conclusion

# The Basics

- ▶ Temporal **L**ogic of **A**ctions
- ▶ Modalities: always ( $\Box$ ) and eventually ( $\Diamond$ )
- ▶ **actions** are logical predicates on **variables**
- ▶ The basic specification “Init and always Next”:

$$Spec \triangleq Init \wedge \Box[Next]_{\langle iter, revlter \rangle}$$

- ▶ Implementation is implication.  $A$  implements  $B$  exactly when  $A \implies B$

## Syntax Quirks

- ▶ Primed variables

$$Next \triangleq iter' = iter + 1 \wedge \text{UNCHANGED } revIter$$

- ▶ Functions

$$list \triangleq [0..42 \mapsto String]$$

- ▶ EXCEPT

$$list' = [list \text{ EXCEPT } ![4] = "hei"]$$

# Table of Contents

Introduction

The Hashmap

TLA+

The Formalization

Results

Lessons Learned and Conclusion

# Overview

Four specifications:

1. Generic hash table
2. Non-concurrent split-order implementing 1.
3. Concurrent split-order
4. (Concurrent with operation ID's)

# Generic Hash Table

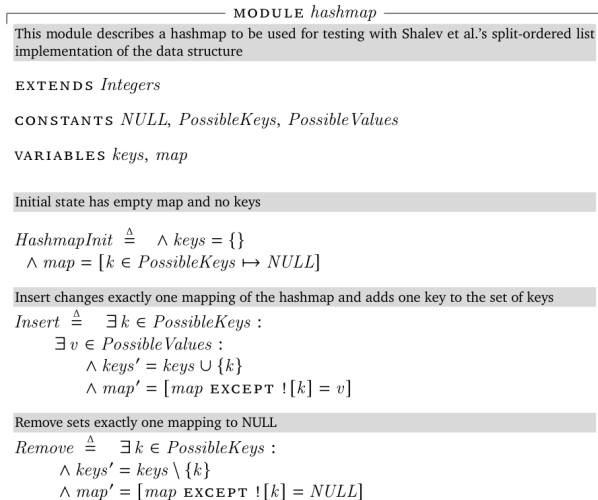


Figure: Generic hash table specification

## Non-concurrent

$SOInsert \triangleq \wedge \exists k \in PossibleKeys :$

$\exists v \in PossibleValues :$

$BucketInsert(k, v)$

$BucketInsert(k, v) \triangleq$

Either a bucket needs to be initialized

$\vee \wedge buckets[k \% size] = NULL$

$\wedge BucketInit(k \% size)$

$\wedge ListInsert(SORRegularKey(k), v)$

$\wedge AuxKeys' = AuxKeys \cup \{k\}$

Or the bucket is already initialized

$\vee \wedge buckets[k \% size] \neq NULL$

$\wedge ListInsert(SORRegularKey(k), v)$

$\wedge AuxKeys' = AuxKeys \cup \{k\}$

$\wedge UNCHANGED \langle buckets \rangle$

$ListInsert(k, v) \triangleq \text{IF } list[k] = NULL$

$\text{THEN } list' = [list \text{ EXCEPT } ![k] = v] \wedge count' = count + 1$

$\text{ELSE } UNCHANGED \langle list, count \rangle$

Figure: The insert operation

# Concurrent Operations I

$$\begin{aligned}
 SONext &\triangleq \\
 &\vee \wedge \exists k \in PossibleKeys : \\
 &\quad \exists v \in PossibleValues : \\
 &\quad \quad Insert(k, v) \\
 &\quad \wedge BagCardinality(activeOps) < MaxActiveOps \\
 &\quad \wedge UNCHANGED \langle buckets, count, list, size \rangle \\
 &\vee \wedge \exists k \in PossibleKeys : \\
 &\quad Delete(k) \\
 &\quad \wedge BagCardinality(activeOps) < MaxActiveOps \\
 &\quad \wedge UNCHANGED \langle buckets, count, list, size \rangle \\
 &\vee Insert1 \\
 &\vee Insert2 \\
 &\vee Insert3 \\
 &\vee Insert4 \\
 &\vee Delete1 \\
 &\vee Delete2 \\
 &\vee Delete3 \\
 &\vee BucketInit1 \\
 &\vee BucketInit2 \\
 &\vee BucketInit3
 \end{aligned}$$

Next step



# Concurrent Operations II

$Insert1 \triangleq$

```

    Start a bucket_init if necessary
 $\exists op \in BagToSet(activeOps) :$ 
     $\wedge op.type = \text{"insert"}$ 
     $\wedge op.step = 1$ 
     $\wedge \text{IF } buckets[op.key \% size] = NULL$ 
        Nextstep and "begin bucket_init" both modify the state of activeOps and need to be combined
        THEN  $activeOps' = (activeOps \ominus SetToBag(\{op\}))$ 
             $\oplus SetToBag(\{[op \text{ EXCEPT } !["step"] = op.step + 1]\})$ 
             $\oplus SetToBag(\{[type \mapsto \text{"bucket_init"}, step \mapsto 1, bucket \mapsto op.key \% size]\})$ 
        ELSE  $NextStep(op)$ 
     $\wedge \text{UNCHANGED } \langle list, buckets, size, count \rangle$ 

```

$Insert2 \triangleq$

```

    If key is already in list, end operation. Else insert in list
 $\exists op \in BagToSet(activeOps) :$ 
     $\wedge op.type = \text{"insert"}$ 
     $\wedge op.step = 2$ 
     $\wedge \text{IF } list[SORegularKey(op.key)] = NULL$ 
        THEN  $list' = [list \text{ EXCEPT } ![SORegularKey(op.key)] = op.value] \wedge NextStep(op)$ 
        ELSE  $End(op) \wedge \text{UNCHANGED } list$ 
     $\wedge \text{UNCHANGED } \langle buckets, size, count \rangle$ 

```

## Concurrent Operations III

$Insert3 \triangleq$

Increment count

$\exists op \in BagToSet(activeOps) :$   
 $\wedge op.type = \text{"insert"}$   
 $\wedge op.step = 3$   
 $\wedge count' = count + 1$   
 $\wedge NextStep(op)$   
 $\wedge UNCHANGED \langle list, buckets, size \rangle,$

$Insert4 \triangleq$

Change size if loadfactor is exceeded and end insert

$\exists op \in BagToSet(activeOps) :$   
 $\wedge op.type = \text{"insert"}$   
 $\wedge op.step = 4$   
 $\wedge \text{IF } count \div size > LoadFactor \text{ THEN } size' = Min(2 * size, MaxSize) \text{ ELSE } UNCHANGED size$   
 $\wedge End(op)$   
 $\wedge UNCHANGED \langle list, buckets, count \rangle$

## Testing the Non-concurrent Spec

$$SOSpec \triangleq SOInit \wedge \Box[SONext]_{\langle keys, AuxKeys, list, buckets, size, count, map \rangle}$$

INSTANCE hashmap

THEOREM  $SOSpec \implies HashmapSpec$

# Testing the Specifications I

Claimed invariants:

1. If a bucket is initialized, then it points to a dummy node in the list

$$\begin{aligned}
 \text{BucketsInitialized} &\triangleq \\
 &\forall i \in 0 \dots (\text{size} - 1) : \\
 &\quad \vee \text{buckets}[i] = \text{NULL} \\
 &\quad \vee \text{buckets}[i] = \text{SODummyKey}(i) \wedge \text{list}[\text{SODummyKey}(i)] = i
 \end{aligned}$$

2. If  $k$  is not in the map, then `delete(k)` fails. Otherwise  $k$  is removed from the map

## Testing the Specifications II

An insert with key in map will not reach step 3

$$\begin{aligned}
 \text{InsertFails} &\triangleq \\
 &\forall op \in \text{OperationStates} : \\
 &\quad \text{IF } op.type = \text{"insert"} \\
 &\quad \text{THEN} \\
 &\quad \quad \wedge op \in \text{BagToSet}(\text{activeOps}) \\
 &\quad \quad \wedge op.step = 1 \\
 &\quad \quad \wedge \text{list}[\text{SORegularKey}(op.key)] \neq \text{NULL} \\
 &\quad \Rightarrow \Box \Diamond (\neg ([op \text{ EXCEPT } !["step"] = 3] \in \text{BagToSet}(\text{activeOps}))) \\
 &\quad \text{ELSE TRUE}
 \end{aligned}$$

3. If  $k$  is in the map, then `insert(k)` fails. Otherwise  $k$  is added to the map

# Table of Contents

Introduction

The Hashmap

TLA+

The Formalization

Results

Lessons Learned and Conclusion

## Non-concurrent

Keys	Values	Diameter	Distinct States	Time (hh:mm:ss)
2	4	10	2,523	00:00:01
2	8	10	23,763	00:00:08
2	16	10	279,075	00:01:02
4	2	17	39,827	00:00:09
4	4	17	1,790,067	00:03:44
4	8	17	147,266,723	07:14:45

Table: Model checking results for SplitOrder

## Concurrent

Invariant	Keys	Values	Conc. Ops	Diameter	Distinct States	Time (mm:ss)
<i>InsertSucceeds</i>	2	2	2	38	10,083	00:02
	2	4	2	38	66,901	00:13
	2	4	3	45	1,351,453	01:50
	4	2	2	62	1,627,390	02:15
<i>DeleteSucceeds</i>	2	2	2	38	10,083	00:01
	2	4	2	38	66,901	00:07
	2	4	3	45	1,351,453	00:31
	4	2	2	62	1,627,390	00:45
<i>BucketsInitialized</i>	2	2	2	38	10,083	00:02
	2	4	2	38	66,901	00:16
	2	4	8	38	624,645	00:16
	2	4	16	38	7,371,157	01:55
	4	2	2	62	1,627,390	00:22
	4	3	2	62	8,368,282	01:39
	4	4	2	62	29,973,646	06:10
	4	5	2	62	85,419,916	18:41
	4	2	3	75	61,353,460	13:59



# Failures

Invariant	Keys	Values	Conc. Ops	Diameter	Distinct States	Time (hh:mm:ss)	Error
Non-Concurrent	4	16		17	84,587,043	02:54:11	3
<i>InsertSucceeds</i>	4	3	2	27	1,426,888	00:33:30	2
<i>DeleteSucceeds</i>	4	3	2	39	6,067,795	00:30:30	1
<i>BucketsInitialized</i>	4	2	4	57	120,175,837	00:26:49	4

Table: Failed model checks

Code	Error
1	GC Overhead limit exceeded
2	Out of heap space
3	No space left on device
4	Unknown

Table: Error Codes

# Table of Contents

Introduction

The Hashmap

TLA+

The Formalization

Results

Lessons Learned and Conclusion

## Designing Specifications

1. Decide on purpose early and precisely
2. Choose step size early and deliberately
3. Test assumed invariants and try to break them

## The Use of Model Checking

- ▶ Poor scaling
- ▶ Computationally expensive
- ▶ Useful
  1. as a way to produce minimal examples of errors, and
  2. as a development tool for algorithms/protocols
- ▶ Increased confidence in algorithms, though arguably no more than a tested implementation

## References

Tables and figures from my bachelor thesis [2]



Ori Shalev and Nir Shavit. “Split-Ordered Lists: Lock-Free Extensible Hash Tables”. In: *J. ACM* 53.3 (2006), pp. 379–405. ISSN: 0004-5411. DOI: [10.1145/1147954.1147958](https://doi.org/10.1145/1147954.1147958). URL: <https://doi.org/10.1145/1147954.1147958>.



Åsmund Aqissiaq Arild Kløvstad. “Formal Verification of a Lock-free Split-order Hashmap”. In: (2020).